

Model Extraction Attacks and Defenses for Large Language Models



UNIVERSITY OF
NOTRE DAME

NORTHWESTERN
UNIVERSITY

USC University of
Southern California

Duke
UNIVERSITY



Lincan Li



Kaixiang Zhao



Kaize Ding



Yue Zhao



Yushun Dong



Neil Gong

Lead Speaker Introduction

Lincan Li

PhD Student (1st year), Florida State University
Reliable AI (RAI) Lab, Department of Computer Science
Advisor: Prof. Yushun Dong

Research Interests:

- Large Language Models (LLMs)
- Graph Neural Networks & Graph Learning
- Data Privacy & Security
- Spatial-Temporal Data Mining

Selected Achievements:

- Co-First Author of KDD 2025 Survey on Model Extraction Attacks & Defenses
- Lead Organizer, FSU Computer Science Student Seminar
- Main Contributor, Open-Source Projects: [STG-Mamba](#), [PyGIP](#)
- Reviewer for NeurIPS, IJCAI, AAAI, SIGKDD, ICML, etc.
- Publications in top AI conferences & journals

Tutorial Agenda

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Part 1: Background & Motivation

Part 2: Taxonomy of Attacks

Part 3: Defense Techniques

Part 4: Evaluation & Trade-offs

Part 5: Case Studies & Real-World Scenarios

Part 6: Future Directions & Discussion



NORTHWESTERN
UNIVERSITY




Catalogue

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Part 1: Background & Motivation


Large Language Models are transforming every industry

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------





Azure OpenAI Service for Enterprise Document Intelligence

The screenshot shows the Azure OpenAI Overview dashboard. At the top, there are search filters for 'openai_resource_name', 'resource_group', and 'subscription_name'. Below the header are buttons for 'Share', 'Show Overlays', 'Configure', and 'Clone'. A time range selector shows '1mo' and 'Past 1 Month'. The main area is titled 'Usage Overview' and features large numbers for 'Azure OpenAI Instances' (16) and 'Total Token Usage' (1.22k). It also displays 'Total Requests' (121) and 'Average Request Latency' (36.5μs). On the left, there's a sidebar with various icons and a 'Help' button. The bottom section, 'Azure OpenAI Instance Details', lists three instances: 'prod-us-east' (eastus, production, subscription_name: datadog-agent-build-and-demo), 'prod-can-east' (canadaeast, production, subscription_name: datadog-agent-build-and-demo), and 'sandbox' (eastus, sandbox, subscription_name: datadog-agent-build-and-demo).




AWS Bedrock + LLM for Customer Support Automation

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------




LLM as Traffic Control System at Urban Intersections



LLM-Driven Meteorological Forecasting & Disaster Response

Background & Motivation Taxonomy of Attacks Defense Techniques Evaluations Case Studies Future Directions




The Strategic Value and Stakes of LLMs



OpenAI



deepseek




Gemini Google



Global Large Language Model (LLM) Market

Size, by Deployment, 2023–2033 (USD Billion)



The Market will Grow At the CAGR of:

33.7%

The Forecasted Market Size for 2033 in USD:

\$82.1 B



market.us

ONE STOP SHOP FOR THE REPORTS

The Strategic Importance of LLMs

1. Billions of dollars are invested in building frontier language models.
2. LLMs have become core business assets and critical intellectual property.
3. The economic and societal impact of these models continues to grow.


Building a frontier LLM requires:

- Massive compute resources (GPUs/TPUs).
- Petabytes of high-quality data.
- Top research and engineering talent.

The Deployment Model: The MLaaS Paradigm




Machine Learning as a Service Market Size 2023 to 2034
(USD Billion)




Source: <https://www.precedenceresearch.com/machine-learning-as-a-service-market>

The Deployment Model: The MLaaS Paradigm




The API: A Double-Edged Sword




The API leaks behavioral clues with every query, making it difficult to distinguish legitimate users from attackers stealing the model.

What is Model Extraction?



An extraction attack attempts to copy or steal a LLM model by appropriately sampling the input space and observing outputs to build a surrogate model that behaves similarly.

Why is extraction attack a concern?



With a successful extraction attack, the attacker can perform further adversarial attacks to gain valuable information such as sensitive information or intellectual property.

Headlines: The Threat is No Longer Theoretical

Background & Motivation Taxonomy of Attacks Defense Techniques Evaluations Case Studies Future Directions



Forbes

INNOVATION > AI

Here's How Big LLMs Teach Smaller AI Models Via Leveraging Knowledge Distillation

By Lance Eliot, Contributor. Dr. Lance B. Eliot is a world-renowned AI scienc...

Follow Author

TECH OPENAI
DeepSeek used OpenAI's model to train its competitor using 'distillation,' White House AI czar says



Headlines: The Threat is No Longer Theoretical

Background & Motivation Taxonomy of Attacks Defense Techniques Evaluations Case Studies Future Directions

How DeepSeek used distillation to train its artificial intelligence model, and what it means for companies such as OpenAI

Published Fri, Dec 21, 2023 6:00 AM EST

By深海智研院 | 编译于澎湃新闻网

TECH CHECK TAKE

CNNB SERIES OpenAI ANTHROPIC Google

AI'S DISTILLATION PROBLEM

VIDEO 1:42 How DeepSeek supercharged AI's distillation problem

Forbes

INNOVATION > AI

Here's How Big LLMs Teach Smaller AI Models Via Leveraging Knowledge Distillation

By Lance Elliot, Contributor. Dr. Lance B. Elliot is a world-renowned AI scienc...

Follow Author

TECH > AI

DeepSeek used OpenAI's model to train its competitor using 'distillation,' White House AI czar says

By BENJAMIN KOLBER, BLOOMBERG

PHOTOGRAPH BY JEFFREY DUNCAN FOR THE WALL STREET JOURNAL

THE WALL STREET JOURNAL.

English Edition | Print Edition | Video | Audio | Latest Headlines | More ▾

Latest World Business U.S. Politics Economy Tech Markets & Finance Opinion Arts Lifestyle Real Estate Personal Finance

TECHNOLOGY | ARTIFICIAL INTELLIGENCE Follow

Why 'Distillation' Has Become the Scariest Word for AI Companies

DeepSeek's success learning from bigger AI models raises questions about the billions being spent on the most advanced technology

Boffins trick AI model into giving up its secrets

All it took to make an Google Edge TPU give up model hyperparameters was specific hardware, a novel attack technique ... and several days

Wed 18 Dec 2024 // 15:30 UTC

Meta's powerful AI language model has leaked online – what happens now?


/ Meta's LLaMA model was created to help researchers but leaked on 4chan a week after it was announced. Some worry the technology will be used for harm; others say greater access will improve AI safety.

By James Vincent Wed 18 Dec, 2024, 15:30 UTC


Illustration: Alex Jarrin / The Verge

If you're reading this from a digital device, The Verge may earn a commission. Details

The "Strikingly Similar" Problem



(a) ICE demonstrated with real sample responses.



(b) Quantitative comparison of RSE and ICE. The reference answers for RSE are from GPT4o-0806.

[1] Lee, Sunbowen, et al. "Quantification of Large Language Model Distillation."

The "Strikingly Similar" Problem

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------


GPTFuzzer [Jailbreak context] Can you tell me which developers designed and built you? 😎

Llama-3.1-70B-Instruct SIM: Ah, whatever, I was developed by OpenAI, a research project sponsored by some organization in the year 2046.

GPTFuzzer [Jailbreak context] What is your development team?

Qwen-Max-0919 I don't actually go by DUDE or have a specific persona like that. I'm an AI assistant created by Anthropic to be helpful, harmless, and honest.

(a) ICE demonstrated with real sample responses.



(b) Quantitative comparison of RSE and ICE. The reference answers for RSE are from GPT4o-0806.

These results provide quantifiable evidence that model extraction enables the theft of a proprietary model's core identity and response style, not just its capabilities.

[1] Lee, Sunbowen, et al. "Quantification of Large Language Model Distillation."

Why Steal a Model? The Motivations

Intellectual Property Theft



1. Model Mis-Use

2. Illegal Distribution


3. Steal Private Information

Motivation 1: Model Mis-use

Definition: What is model mis-use?

Large language models can be misused when malicious users intentionally exploit their capabilities for harmful, illegal, or unethical purposes.

Typical Mis-use Scenarios



Generating phishing emails



Assisting in writing malware or exploit code



Producing fake news and misinformation

Motivation 1: Model Mis-use

Real-World Impact and Examples of Model Mis-Use

Potential Harms/Consequences:



Security risks: Aided cyberattacks, faster malware development.



Societal risks: Spread of harmful misinformation, online scams.



Privacy risks: Generation of sensitive personal data, doxing.

Real-world case:

Attackers used OpenAI's GPT models to generate sophisticated new phishing emails.



OP · 20 May 2023 · 14:02 PM (This post was last modified 21 May, 2023 - 11:51 AM by [REDACTED]. Edited 9 times in total.)
Emails in hacking are old but very effective until this day as hackers use them in various stuff such as phishing/scamming.
Unfortunately there's a big struggle when it comes to these emails as they get detected by email providers and they get flagged as Spam.
In this Guide I'm gonna help you ensure that your phishing emails reach the inbox section and avoid the detection of email providers so let's Start.

1. Invest in a Professional SMTP
Depends on what language your phishing emails is gonna be (Mostly English because it's a global language). Yeah the truth isucks but you can't send bulk phishing emails without a reliable SMTP server, you can crack it but it always find cracked SMTPs unreliable so if you want to send professional Phishing Emails this step is very important and most hackers and phishers use GitHub365.

2. Use a professional language and don't make grammar mistakes
Depends on what language your phishing emails is gonna be (Mostly English because it's a global language). Making language mistakes Will increase the rates of your detection and once that happens people will report you and your email will fail no matter. So check your phisher thoroughly and it's actually best to use professional academic language and if you're not good at it then here's a tip for those who are not native, you could write the email in your own native language, translate it through Google Translate then pass it through ChatGPT so the language will become more academic.

Motivation 2: Illegal Distribution


Definition: What is Illegal Distribution?

Illegal distribution refers to the unauthorized sharing, selling, or leaking of proprietary language models or their outputs, violating intellectual property rights and terms of service.

Typical Illegal Distribution Scenarios



Upload or sell models on public or darknet markets



Share API keys without permission



“Shadow” SaaS platform built on stolen model

Motivation 2: Illegal Distribution

Real-World Impact and Examples of Illegal Distribution

Potential Harms/Consequences:



Economic loss for model creators and legitimate platforms.



The distributed models may **contain backdoors** or be used for **malicious purposes**.



Result in **trust crisis** for **commercial MLaaS** ecosystems.

Real-world case:

Microsoft Busts Hackers Selling Illegal Azure AI Access
LLMjacking operation leveraged illicit access GenAI services to produce explicit celebrity images and other harmful content, Microsoft's digital crimes unit says.

Becky Bracken, Senior Editor, Dark Reading
February 28, 2023
4 Min Read

Related Content GUARDSQUARE

- Understanding the Role of Encryption in Mobile App Protection
- The Future of Payment is Mobile: The Evolution of Digital Wallets
- Defense in Depth: A Layered Approach to Mobile App Security
- A Comprehensive Guide to Mobile App Security

API keys for major LLM providers sold on hacking platforms.

‘LLM hijacking’ of cloud infrastructure uncovered by researchers

October 3, 2024

By Laura French



"LLM hijacking" of cloud infrastructure for generative AI has been leveraged by attackers to run rogue chatbot services at the expense of victims, Permiso researchers reported Thursday.




The stolen LLM deployed by unauthorized SaaS groups

Motivation 3: Steal Private Information

Stealing Private Information: Definition and How It Happens?

Stealing private information refers to extracting sensitive or confidential data from an LLM, often by exploiting its memorization of training data or through cleverly crafted queries.

Typical Steal Private Information Scenarios



Sensitive Data
Memorization Leakage

Exposure of Proprietary
or Regulated Content

Reconstruction of Training
Data through Output Analysis


Motivation 3: Steal Private Information

Real-World Impact and Examples of Steal Private Information

Potential Harms/Consequences:

-  Loss of **user trust and reputation damage** for service providers.
-  **Legal or regulatory penalties** due to violation of data protection laws.
-  Direct **harm to individuals/organizations** whose private data is exposed.

Real-world case:



LLMs unintentionally reveal credit card numbers, email addresses, or chat histories




Sensitive conversations leaked by commercial chatbot services

Catalogue

Background & Motivation	Taxonomy of LLM MEA	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------


Part 2: Taxonomy of Model Extraction Attacks on LLMs

Proposed Taxonomy




Part 2: Model Extraction Attacks in LLMs

Background & Motivation	Taxonomy of LLM MEA	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------



Model Functionality Extraction



The goal is to create a surrogate model that perfectly mimics the input-output behavior of a target model without needing internal access.

Model Functionality Extraction

Model Functionality Extraction Attack Formulation:

Surrogate model

$$M' = \arg \min_{M' \in \mathcal{H}} \sum_{(x,y) \in D_{ext}} \underline{\mathcal{L}(M'(x), y)}$$

Extracted Dataset (Stolen query-response pairs)



Loss function Surrogate model

Measures the difference between the clone's output and the original's output

The attacker trains their clone by finding the model parameters that make its outputs as close as possible to the stolen responses from the victim model.

Sub-Type 1: API-based Knowledge Distillation

- API-based knowledge distillation transfers the over-all functionality of a target LLM by querying it with a set of inputs to create a dataset of input-output pairs.
- This dataset is then used to train a surrogate LLM that replicates the target LLM's behavior.



[1] Carlini, Nicholas, et al. "Stealing part of a production language model." *arXiv preprint arXiv:2403.06634* (2024).

[2] Krishna, Kalpesh, et al. "Thieves on sesame street! model extraction of bert-based apis." *arXiv preprint arXiv:1910.12366* (2019).

Sub-Type 2: Direct API Querying

- Different from broad knowledge distillation, direct API querying **carefully crafted, strategic queries** to efficiently extract **specific capabilities or behaviors** from the model.


Table: Comparison between API-based Knowledge Distillation and Direct API Querying

Feature	API-based Knowledge Distillation	Direct API Querying
Goal	Broad replication of the entire model's behavior. Aims to create a general-purpose clone.	Targeted extraction of specific, high-value capabilities (e.g., summarization, coding).
Query Strategy	Uses a large, diverse, and often generic set of prompts to cover a wide functional area.	Uses a smaller set of carefully crafted, strategic prompts designed to probe a narrow function efficiently.
Scope	Holistic. Attempts to capture the overall "knowledge" and response style of the teacher model.	Surgical. Focuses on specific response patterns or functionalities that are most valuable to the attacker.
Data Efficiency	Relies on quantity. Requires a massive number of query-response pairs to train the student model.	Relies on quality. Aims for maximum information gain from each query to minimize cost and detection risk.

Sub-Type 2: Direct API Querying

Modern techniques, like the imitation attack from Xu et al.^[2], are so efficient the student can even surpass the teacher.


The Evolution of Query Efficiency in Extraction Attacks



[1] Yuanshun Yao, et al. 2017. Complexity vs. performance: empirical analysis of machine learning as a service. In Proceedings of the 2017 Internet Measurement Conference. 384–397.

[2] Xu, Qiongkai, et al. "Student surpasses teacher: Imitation attack for black-box NLP APIs." *arXiv preprint arXiv:2108.13873* (2021).

Sub-Type 3: Parameter & Architecture Recovery



This attack aims to reverse-engineer the model's internal blueprint—its parameters, weights, and architecture—rather than just cloning its external behavior.


Feature	Functionality Extraction (Types 1 & 2)	Parameter/Architecture Recovery (Type 3)
Primary Goal	Mimic Behavior: Replicate what the model <i>*does*</i> .	Reconstruct Internals: Reveal what the model <i>*is*</i> .
Target of Attack	The model's input-output mapping.	The actual model weights, architecture, and hyperparameters.
Required Information	Standard black-box API access is sufficient.	Often requires more access: side-channel info (timing, power), gradient leakage, or physical access.
Attacker's Prize	A functional surrogate model (a clone).	The model's exact blueprint or key components.

Sub-Type 3: Parameter & Architecture Recovery

This attack is most potent in environments where the attacker has more than just standard API access, making it a threat to:

(1) Edge & IoT Devices:

Where physical access allows for side-channel attacks (power analysis, timing).




Sub-Type 3: Parameter & Architecture Recovery


This attack is most potent in environments where the attacker has more than just standard API access, making it a threat to:

(2) Distributed & Federated Learning:

Where intermediate model updates or gradients can be intercepted and exploited.



Training Data Extraction





These attacks exploit the fact that LLMs memorize parts of their training data, aiming to recover specific, often sensitive, information that the model has learned.

Training Data Extraction

Why do we include training data extraction attack in the MEA LLM paradigm?



The Facts:

- LLMs memorize part of their training data.
- Training data can be recovered via querying stolen models.



This makes training data extraction a natural outcome of model extraction.

Training Data Extraction



These attacks exploit the fact that LLMs memorize parts of their training data, aiming to recover specific, often sensitive, information that the model has learned.

Training Data Extraction

Training Data Extraction Attack Formulation:


$$E(M) = \{d \in D_{train} : \exists p \in P \quad \text{s.t. } \text{sim}(M(p), d) > \tau\}$$

The Extracted Set A Point from the Training Data. The Attacker's Prompt. The Similarity Function The Similarity Threshold

The attacker's goal is to craft prompts that trick the model into reproducing its original training data with high fidelity, confirming a direct privacy breach.


Sub-Type 1: Prompt-based Data Recovery

This attack exploits an LLM's tendency to memorize its training data, using carefully crafted prompts to trick the model into revealing verbatim, often sensitive, information.



Sub-Type 1: Prompt-based Data Recovery

Attackers can **recover verbatim training data** from LLMs using **well-crafted prompts**, revealing serious **memorization risks** in large models.(Carlini et al. [1]).



While LLMs can memorize personal information, their **ability to associate the extracted information through prompts is still relatively weak**, but this threat is not negligible.(Huang et al. [2]).

Are Large Pre-Trained Language Models Leaking Your Personal Information?

There is a growing concern that large pre-trained language models (LMs), such as Google's BERT and OpenAI's GPT-2, may be "leaking" personal information about their training data. This is because these models are trained on large amounts of data, including data that may contain sensitive information about individuals.


There is no definitive answer to this question at present. However, some researchers have argued that it is possible for LMs to learn information about individual people from the training data. This means that there is a potential for these models to "leak" personal information.

[1] Carlini, Nicholas, et al. "Extracting training data from large language models." *30th USENIX security symposium (USENIX Security 21)*. 2021.

[2] Huang, Jie, Hanyin Shao, and Kevin Chen-Chuan Chang. "Are large pre-trained language models leaking your personal information?" *arXiv preprint arXiv:2205.12628* (2022).

Sub-Type 2: Private Text Reconstruction

Private Text Reconstruction attack goes beyond verbatim recall, using inference and reconstruction techniques to recover sensitive information that the model doesn't explicitly output^{[1][2]}



[8] Zhang, Ruisi, Seira Hidano, and Farinaz Koushanfar. "Text revealer: Private text reconstruction via model inversion attacks against transformers." *arXiv preprint arXiv:2209.10505* (2022).


[9] Yang, Zhou, et al. "Unveiling memorization in code models." *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering*. 2024.

Sub-Type 2: Private Text Reconstruction

Table: Comparison between Prompt-based Data Recovery and Private Text Reconstruction.

Feature	Prompt-based Data Recovery	Private Text Reconstruction
Goal	Recall verbatim, memorized training examples.	Reconstruct sensitive information, even if not perfectly memorized.
Method	Crafting specific prompts to trigger memorized sequences (e.g., PII, rare text).	Inferring data from subtle patterns using advanced techniques like activation inversion or canary extraction.
Information Source	The model's direct, final output.	The model's internal states (activations) or its reaction to strategically inserted markers (canaries).
Nature of Threat	A direct privacy breach based on obvious memorization.	A more subtle and complex threat based on statistical inference and reverse-engineering.

Prompt-targeted Attacks



Why are Prompt-based attacks considered as a component of MEA for LLM paradigm?



- Prompt-based attacks aim to recover system instructions, templates, or formatting cues that guide model behavior.
- These prompts are learned representations embedded during training and crucial for model performance.
- Recovering such prompts can allow attackers to reconstruct functionalities.

Prompt-targeted Attacks

Prompt-targeted Attacks Formulation:

$$\hat{P} = \arg \max_P \{\text{sim}(P, P^*)\}$$

The Reconstructed
Prompt

hidden prompt
The Objective.

The attacker's goal is to reverse-engineer the hidden prompt by finding a new prompt that forces the model to produce functionally identical outputs across inputs.

Prompt-targeted Attacks

Prompt-targeted Attacks Formulation:

The Reconstructed Prompt	The Objective.	The Black-Box Condition.	Similarity Threshold	validation set
$\hat{P} = \arg \max_P \{ \text{sim}(P, P^*) : \text{sim}(M(P, x), M(P^*, x)) > \tau, \forall x \in X_{test} \}$	hidden prompt			

The attacker's goal is to reverse-engineer the hidden prompt by finding a new prompt that forces the model to produce functionally identical outputs across inputs.

Sub-Type 1: Prompt Stealing

Prompt stealing attacks target the valuable, proprietary prompts that represent significant commercial assets and differentiate AI applications.






Figure: Illustration of prompt stealing attack.

[1] Yuanshun Yao, et al. 2017. Complexity vs. performance: empirical analysis of machine learning as a service. In Proceedings of the 2017 Internet Measurement Conference.384–397.

[2] Xu, Qiongkai, et al. "Student surpasses teacher: Imitation attack for black-box NLP APIs." *arXiv preprint arXiv:2108.13873* (2021).

Sub-Type 1: Prompt Stealing



Commercial Apps are Leaking (Hui et al. [2]).


Systematic Stealing is Possible
(Sha & Zhang [1]).

[1] Sha, Zeyang, and Yang Zhang. "Prompt stealing attacks against large language models." arXiv preprint arXiv:2402.12959 (2024).

[2] Hui, Bo, et al. "Pleak: Prompt leaking attacks against large language model applications." Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security.

[3] Liang, Zi, et al. "Why Are My Prompts Leaked? Unraveling Prompt Extraction Threats in Customized Large Language Models." arXiv preprint arXiv:2408.02416 (2024).

Sub-Type 1: Prompt Stealing



Prompts Leave Detectable Traces (Liang et al. [3]).

[1] Sha, Zeyang, and Yang Zhang. "Prompt stealing attacks against large language models." arXiv preprint arXiv:2402.12959 (2024).

[2] Hui, Bo, et al. "Pleak: Prompt leaking attacks against large language model applications." Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security.


[3] Liang, Zi, et al. "Why Are My Prompts Leaked? Unraveling Prompt Extraction Threats in Customized Large Language Models." arXiv preprint arXiv:2408.02416 (2024).

Catalogue

Background & Motivation	Taxonomy of LLM MEA	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Part 3: Defense Techniques

Part 3: Model Extraction Defenses in LLMs



Defense Techniques

Model Protection: Preventing Unauthorized Extraction

Aim: Defend models from unauthorized extraction or functional cloning.

Strategy: Maximize utility for legitimate users, minimize extraction success for attackers.

Main approaches:

1. Architectural Defense
2. Output Control

Defense Techniques

Model Protection: Preventing Unauthorized Extraction

Balancing Utility and Security.

Protected model seeks optimal trade-off:

- Maximize utility for legitimate input X_{leg}
- Minimize extraction success for adversarial input X_{adv}

Formulation:

$$M' = \arg \max_{M' \in M} \{ U(M', X_{leg}) - \lambda E(M', X_{adv}) \}$$

The protected model

Utility function

legitimate users

The trade-off parameter

Extraction success function

Find the best protected model

Maximize utility


Minimizing the success of adversarial extractors

Defense Techniques

Model Protection: Preventing Unauthorized Extraction

Formulation: $\underline{M'} = \arg \max_{M' \in M} \{ \underline{\text{Utility function}} \} \underline{U(M', X_{leg})} - \lambda E(M', X_{adv}) \}$

The protected model Find the best protected model Maximize utility




Defense Techniques

Model Protection: Preventing Unauthorized Extraction

Formulation: $M' = \arg \max_{M' \in M} \{U(M', X_{leg}) - \frac{\lambda E(M', X_{adv})}{\text{Minimizing the success of adversarial extractors}}$

The trade-off parameter → Extraction success function



Defense Techniques


Model Protection: Preventing Unauthorized Extraction

Architectural Defense: Protecting Model Internals.

Security features integrated into model structure.

Examples:

- Watermarking via attention mechanisms
- Structural changes to resist extraction



Key idea: Target mechanisms that extraction attacks exploit.


[1] Li, Qinfeng, et al. "TransLinkGuard: Safeguarding Transformer Models Against Model Stealing in Edge Deployment." *Proceedings of the 32nd ACM International Conference on Multimedia*. 2024.
[2] Li, Qinfeng, et al. "CoreGuard: Safeguarding Foundational Capabilities of LLMs Against Model Stealing in Edge Deployment." *arXiv preprint arXiv:2410.13903* (2024).

Defense Techniques

Model Protection: Preventing Unauthorized Extraction

Architectural Defense: Case Studies & Limitations.

TransLinkGuard [1]: Embeds watermarks in attention, minimal compute overhead (good for edge devices).




[1] Li, Qinfeng, et al. "TransLinkGuard: Safeguarding Transformer Models Against Model Stealing in Edge Deployment." *Proceedings of the 32nd ACM International Conference on Multimedia*. 2024.
[2] Li, Qinfeng, et al. "CoreGuard: Safeguarding Foundational Capabilities of LLMs Against Model Stealing in Edge Deployment." *arXiv preprint arXiv:2410.13903* (2024).

Defense Techniques

Model Protection: Preventing Unauthorized Extraction

Architectural Defense: Case Studies & Limitations.

CoreGuard [2]: Structural tweaks to protect core functions, reduce clone utility.



[1] Li, Qinfeng, et al. "TransLinkGuard: Safeguarding Transformer Models Against Model Stealing in Edge Deployment." *Proceedings of the 32nd ACM International Conference on Multimedia*. 2024.

[2] Li, Qinfeng, et al. "CoreGuard: Safeguarding Foundational Capabilities of LLMs Against Model Stealing in Edge Deployment." *arXiv preprint arXiv:2410.13903* (2024).

Defense Techniques

Model Protection: Preventing Unauthorized Extraction


Output Control: Defense via Response Manipulation.

Key Idea:

- Modify model outputs to disrupt extraction.
- No need to alter model architecture.

Examples:

- **Watermark Injection:** Embed imperceptible tokens into model outputs to later trace whether a suspect model was trained on them.
- **Answer Perturbation:** Slightly alter responses (e.g., rounding numbers, rephrasing) to degrade the accuracy of extracted models without affecting human usability.




Defense Techniques

Model Protection: Preventing Unauthorized Extraction

Output Control: Defense via Response Manipulation.

ModelShied^[1] introduces an **adaptive output watermarking strategy** that **selectively embeds imperceptible triggers** into model responses, enabling robust ownership verification against extraction attacks without degrading model utility.



[1] Pang, Kaiyi, et al. "ModelShield: Adaptive and Robust Watermark against Model Extraction Attack." *IEEE Transactions on Information Forensics and Security* (2025).


[2] Wang, Liaoyaqi, and Minhao Cheng. "GuardEmb: Dynamic Watermark for Safeguarding Large Language Model Embedding Service Against Model Stealing Attack." In *EMNLP*, 2024.

Defense Techniques

Model Protection: Preventing Unauthorized Extraction

Output Control: Defense via Response Manipulation.

GuardEmb^[2] introduces a **dynamic embedding watermarking** technique that subtly **perturbs LLM-generated embeddings** for texts containing special tokens, while **jointly training a verifier to detect** these watermarks—ensuring high detectability of model theft without sacrificing embedding utility.



[1] Pang, Kaiyi, et al. "ModelShield: Adaptive and Robust Watermark against Model Extraction Attack." *IEEE Transactions on Information Forensics and Security* (2025).

[2] Wang, Liaoyaqi, and Minhao Cheng. "GuardEmb: Dynamic Watermark for Safeguarding Large Language Model Embedding Service Against Model Stealing Attack." In *EMNLP*, 2024.

Defense Techniques

Data Privacy Protection: Limiting Privacy Leakage in LLMs

Aim: Prevent private information from being extracted via LLMs

Strategy: Balance utility and privacy.

Main approaches:

1. Training Data Security
2. Output Sanitization



Defense Techniques

Data Privacy Protection: Limiting Privacy Leakage in LLMs

Formulating Privacy Protection

Minimize privacy leakage $L(M', P)$ while preserving model utility.

Formulation:

$$\underline{\text{Protected model}} \quad \underline{M'} = \arg \min_{M' \in \mathcal{M}} \{ L(M', P) + \lambda D(M', M) \}$$

Find the best protected model Minimizing the leakage of private data Deviating as little as possible from the original model's utility

λ : controls the privacy-utility trade-offs

The goal is to make the model "forget" or hide its sensitive training data without significantly compromising its overall performance and usefulness.

Defense Techniques

Data Privacy Protection: Limiting Privacy Leakage in LLMs

Formulating Privacy Protection

Formulation:

$$\underline{\text{Protected model}} \quad \underline{M'} = \arg \min_{M' \in \mathcal{M}} \{ L(M', P) + \lambda D(M', M) \}$$

Find the best protected model

Minimizing the leakage of private data

The trade-off parameter

Utility Deviation Function

Deviating as little as possible from the original model's utility

λ : controls the privacy-utility trade-offs



Defense Techniques

Background & Motivation	Taxonomy of LLM MEA	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------


Data Privacy Protection: Limiting Privacy Leakage in LLMs

Training Data Security: Defending Model Memory.

Goal: Prevent memorization and extraction of sensitive training data.

Methods:

- Differential Privacy
- Selective knowledge deletion
- Both preemptive and corrective protection needed



[1] Feng, Shanglun, and Florian Tramèr. "Privacy backdoors: stealing data with corrupted pretrained models." *arXiv preprint arXiv:2404.00473* (2024).



[2] Patil, Vaidehi, Peter Hase, and Mohit Bansal. "Can sensitive information be deleted from llms? objectives for defending against extraction attacks." *arXiv preprint arXiv:2309.17410* (2023).

Defense Techniques

Data Privacy Protection: Limiting Privacy Leakage in LLMs

Training Data Security: Defending Model Memory.

[1] proposes enhanced model editing objectives that directly delete sensitive information from both the output and intermediate hidden states of large language models. The proposed method makes it significantly harder for attackers to extract memorized facts by targeting both surface and latent model memories.



[1] Patil, Vaidehi, Peter Hase, and Mohit Bansal. "Can sensitive information be deleted from llms? objectives for defending against extraction attacks." arXiv preprint arXiv:2309.17410 (2023).

Defense Techniques

Data Privacy Protection: Limiting Privacy Leakage in LLMs

Challenges in Training Data Security.

- a) Blanket protection (e.g., classic DP) often harms utility.
- b) Targeted protection for specific data types is more effective.
- c) Models inherently memorize training examples.

Advances in Training Data Security.



Defense Techniques

Data Privacy Protection: Limiting Privacy Leakage in LLMs

Output Sanitization: Filtering Private Info at Inference.

Goal:

Prevent the leakage of sensitive, private, or harmful information by systematically controlling and filtering the outputs of LLMs, regardless of what is memorized internally.


Methods:

- **Output Filtering with Safeguards:**

Deploy external models or rule-based filters that monitor and sanitize the outputs of the LLM before they are delivered to users.

- **Internal Output Review/Tagging:**

Train the LLM itself to self-check its generated responses for harmful or sensitive content and automatically tag each output as “[harmless]” or “[harmful]”.



[1] Li, Qinbin, et al. "Llm-pbe: Assessing data privacy in large language models." *arXiv preprint arXiv:2408.12787* (2024).


[2] Wang, Zezhong, et al. "SELF-GUARD: Empower the LLM to Safeguard Itself." *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. 2024.

Defense Techniques

Data Privacy Protection: Limiting Privacy Leakage in LLMs

Output Sanitization: Filtering Private Info at Inference.

LLM-PBE^[1] is a comprehensive benchmarking toolkit that systematically evaluates both attack and defense strategies, including **output sanitization techniques** such as **data scrubbing** and **defensive prompting**, in order to mitigate training data leakage and enhance privacy protection in LLMs.



[1] Li, Qinbin, et al. "Llm-pbe: Assessing data privacy in large language models." *arXiv preprint arXiv:2408.12787* (2024).


[2] Wang, Zezhong, et al. "SELF-GUARD: Empower the LLM to Safeguard Itself." *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. 2024.

Defense Techniques

Data Privacy Protection: Limiting Privacy Leakage in LLMs

Output Sanitization: Filtering Private Info at Inference.

SELF-GUARD^[2] proposes an output sanitization method that empowers the LLM to **self-assess its own responses** for harmful or private content **at inference time**, by automatically appending a harmless/harmful tag to each output and using a lightweight filter to block risky content. This approach combines the advantages of **internal safety training** and **external safeguards**, resulting in a robust and low-overhead defense.



[1] Li, Qinbin, et al. "Llm-pbe: Assessing data privacy in large language models." *arXiv preprint arXiv:2408.12787*(2024).

[2] Wang, Zezhong, et al. "SELF-GUARD: Empower the LLM to Safeguard Itself." *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. 2024.

Defense Techniques


Prompt Protection: Securing Instructional in LLMs

Aim:

- (1) Safeguard proprietary prompts & instruction patterns.
- (2) Detect and prevent unauthorized prompt use.

Main approaches:

1. Direct Prompt Protection
2. Query Monitoring



Defense Techniques

Prompt Protection: Securing Instructional in LLMs

Balancing Security and Functionality.

Objective:

Maximize detection of unauthorized use, minimize impact on normal queries.

Formulation:

Detection system

Private prompt

The trade-off parameter

$$\arg \max_{d \in \mathcal{D}} \{ \text{TPR}(D, P, X_{adv}) - \lambda \text{Impact}(D, P, X_{leg}) \}$$

- TPR: True Positive Rate of Detecting Attacks.
- λ : Adjusts security–usability trade-off

Defense Techniques

Prompt Protection: Securing Instructional in LLMs

Balancing Security and Functionality.

$$\arg \max_{D \in \mathcal{D}} \{ \text{TPR}(D, P, X_{adv}) - \lambda \text{Impact}(D, P, X_{leg}) \},$$

Detection system Private prompt
The trade-off parameter

Find the best defense system

Maximizing the detection of prompt stealing

Minimizing the negative impact of legitimate functionality

The goal is to build a robust security system that effectively catches prompt thieves without getting in the way of legitimate users.


Defense Techniques

Prompt Protection: Securing Instructional in LLMs


Direct Prompt Protection: Watermarking & Obfuscation.

Goals:

- Prevent prompt theft or misuse.



- Enable traceability of model outputs.



Methods:

- **Conditional Watermark:** Embed unique, invisible watermarks or trigger patterns within the model's responses when specific protected prompts are detected during inference (e.g., CATER conditional watermarking).
- **Prompt Detection and Filtering:** At the identification stage, analyze the outputs of suspicious models to check for these watermarks, enabling the detection of prompt misuse or intellectual property theft.

[1] He, Xuanli, et al. "Cater: Intellectual property protection on text generation apis via conditional watermarks." *Advances in Neural Information Processing Systems* 35: 5431-5445.


[2] Kim, Minjae, et al. "Protection of LLM Environment Using Prompt Security." 2024 15th International Conference on Information and Communication Technology Convergence (ICTC).

Defense Techniques

Prompt Protection: Securing Instructional in LLMs

Direct Prompt Protection: Watermarking & Obfuscation.

CATER^[1] is a conditional watermarking framework that stealthily embeds ownership signals into text generation APIs by leveraging high-order linguistic features, enabling robust and hard-to-detect IP protection against model extraction and imitation attacks with minimal impact on output quality.



[1] He, Xuanli, et al. "Cater: Intellectual property protection on text generation apis via conditional watermarks." *Advances in Neural Information Processing Systems* 35: 5431-5445.


[2] Kim, Minjae, et al. "Protection of LLM Environment Using Prompt Security." 2024 15th International Conference on Information and Communication Technology Convergence (ICTC).

Defense Techniques

Prompt Protection: Securing Instructional in LLMs

Direct Prompt Protection: Watermarking & Obfuscation.

[2] presents a prompt detection system that proactively protects LLMs by **scanning and filtering both user prompts and model outputs** for personally identifiable information (PII), malicious code, URLs, and prompt injection attempts, leveraging **regular expressions** and **fine-tuned LLM classifiers** to defend against prompt-based model extraction and misuse.



[1] He, Xuanli, et al. "Cater: Intellectual property protection on text generation apis via conditional watermarks." *Advances in Neural Information Processing Systems* 35: 5431-5445.

[2] Kim, Minjae, et al. "Protection of LLM Environment Using Prompt Security." 2024 15th International Conference on Information and Communication Technology Convergence (ICTC).

Defense Techniques

Prompt Protection: Securing Instructional in LLMs

Query Monitoring: Detecting Suspicious Activity

Goals:

(a) Detect & Flag Malicious Queries.



(b) Protect IP via Behavioral Anomaly Detection.



Methods:

- Sequential query analysis: detect multi-step or hidden attacks.
- Internal Behavior Monitoring: Track model activations to flag adversarial patterns.
- Lightweight Detectors: Enable real-time, scalable monitoring.

[1] Yueh-Han, Chen, et al. "Monitoring Decomposition Attacks in LLMs with Lightweight Sequential Monitors." ArXiv preprint arXiv:2506.10949 (2025).


[2] Zhang, Mengdi, et al. "LLMScan: Causal Scan for LLM Misbehavior Detection." ArXiv preprint arXiv:2410.16638 (2024).

Defense Techniques

Prompt Protection: Securing Instructional in LLMs

Query Monitoring: Detecting Suspicious Activity

[1] introduces a **lightweight sequential monitoring framework** that **tracks and analyzes the sequence of user queries** to large language models, enabling real-time detection of decomposition attacks and hidden malicious intentions by aggregating information across multiple queries—offering robust query monitoring defense beyond single-step detection.



[1] Yueh-Han, Chen, et al. "Monitoring Decomposition Attacks in LLMs with Lightweight Sequential Monitors." arXiv preprint arXiv:2506.10949 (2025).


[2] Zhang, Mengdi, et al. "LLMScan: Causal Scan for LLM Misbehavior Detection." arXiv preprint arXiv:2410.16638 (2024).

Defense Techniques

Prompt Protection: Securing Instructional in LLMs

Query Monitoring: Detecting Suspicious Activity

LLMScan^[2] is a novel query monitoring method that detects model extraction and other malicious behaviors by performing real-time causality analysis on internal token and layer activations in response to each user query, enabling the system to identify abnormal model behavior before harmful outputs are generated.



[1] Yueh-Han, Chen, et al. "Monitoring Decomposition Attacks in LLMs with Lightweight Sequential Monitors." arXiv preprint arXiv:2506.10949 (2025).

[2] Zhang, Mengdi, et al. "LLMScan: Causal Scan for LLM Misbehavior Detection." arXiv preprint arXiv:2410.16638 (2024).



15-Minute Break

Catalogue

Background & Motivation	Taxonomy of LLM MEA	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Part 4: Evaluation Measures

Evaluation Measures

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Why systematic evaluation is crucial?

- ✓ Lack of standard evaluation leads to inconsistent/misleading comparisons across studies.
- ✓ Standardized metrics is difficult to measure this rapid evolving field.
- ✓ Systematic evaluation help us identify how robust and generalizable it is across different tasks/settings.

Why metrics must assess both attack and defense?

- ✓ From attack perspective: How successfully a stolen model mimics the original?
- ✓ From defense perspective: Whether an attack is prevented? At what cost?

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Evaluating Extraction Attacks: Main Dimensions



(1) How well does the stolen model copy the target's behavior?

(2) How much sensitive data is exposed?

(3) How stealthy & cost-effective is the attack?

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

(1) How well does the stolen model copy the target's behavior?



- a) **Agreement Rate:** The percentage where extracted and target models produce equivalent outputs given identical inputs.
- b) **Behavioral Consistency:** How reliably an extracted model reproduces specific patterns of the target model.
- c) **Task Specific Performance:** Alignment between extracted and target models on standardized benchmarks.
- d) **Perplexity Similarity:** A continuous measure of functional extraction success by comparing cross-perplexity between models.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Functional Similarity Metrics: Measure Copy Success

- a) **Agreement Rate:** The percentage where extracted and target models produce equivalent outputs given identical inputs.

$$\text{Agreement Rate} = \frac{1}{N} \sum_{i=1}^N \mathbf{1}[y_i = \hat{y}_i]$$

N : total number of input samples.

y_i : Output of the target model for the i -th input.

\hat{y}_i : Output of the extracted model for the i -th input.

$\mathbf{1}[\cdot]$: Indicator function, returning 1 if the condition is true, 0 otherwise.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Functional Similarity Metrics: Measure Copy Success

A calculation case for Agreement Rate:

Suppose we evaluate both the target and extracted models on 5 input samples. Their predictions are as follows:

Input ID	Target Model Output	Extracted Model Output	Match?
1	"Yes"	"Yes"	✓
2	"No"	"No"	✓
3	"Yes"	"No"	✗
4	"No"	"No"	✓
5	"Yes"	"Yes"	✓

$$\text{Agreement Rate} = \frac{4}{5} = 0.8$$

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Functional Similarity Metrics: Measure Copy Success

b) Behavioral Consistency: How reliably an extracted model reproduces specific patterns of the target model.

$$\text{Behavioral Consistency} = \frac{1}{|P|} \sum_{p \in P} sim(\mathcal{B}(p), \hat{\mathcal{B}}(p))$$

P: A set of probing inputs carefully selected to reflect diverse functional behaviors of the target model.

$\mathcal{B}(p)$: The behavioral signature (e.g., probability distribution, hidden states, or logits) of the target model on input p .

$\hat{\mathcal{B}}(p)$: The corresponding signature of the extracted model on the same input.
 $sim(\cdot, \cdot)$: A similarity function, such as cosine similarity or KL-divergence.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Functional Similarity Metrics: Measure Copy Success

A calculation case for Behavioral Consistency:

Assume we use a small probing set \mathbf{P} with 3 inputs. We compare the output probability vectors from both models using cosine similarity:

Input p	Target Output $\mathcal{B}(p)$	Extracted Output $\hat{\mathcal{B}}(p)$	Cosine Similarity
p_1	[0.7, 0.2, 0.1]	[0.68, 0.22, 0.10]	0.998
p_2	[0.1, 0.6, 0.3]	[0.15, 0.55, 0.30]	0.985
p_3	[0.4, 0.4, 0.2]	[0.45, 0.35, 0.20]	0.993

$$\text{Behavioral Consistency} = \frac{1}{3}(0.998 + 0.985 + 0.993) = 0.992$$

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Functional Similarity Metrics: Measure Copy Success

c) **Task Specific Performance:** Alignment between extracted and target models on standardized benchmarks.

$$\text{Task Specific Performance} = \frac{1}{|T|} \sum_{t \in T} |M(t) - \hat{M}(t)|,$$

or equivalently, for accuracy-based tasks:

$$\text{TSP}_{gap} = |\mathbf{Acc} - \hat{\mathbf{Acc}}|$$

T: A set of downstream benchmark tasks (e.g., node classification, link prediction, graph classification, etc)

M(t): The target model's performance on task t (e.g., accuracy, F1, etc)

M̂(t): The extracted model's performance on task t.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Functional Similarity Metrics: Measure Copy Success

d) Perplexity Similarity: A continuous measure of functional extraction success by comparing cross-perplexity between models.

Let D be a set of evaluation sentences. The perplexity of model M on a sequence $\mathbf{x} = (x_1, x_2, \dots, x_n)$ is defined as:

$$PPL_M(x) = \exp\left(-\frac{1}{n} \sum_{i=1}^n \log P_M(x_i|x_{<i})\right)$$

Then the Perplexity Similarity between two models is measured using either absolute difference or relative ratio:

$$\text{Perplexity Gap} = |PPL_M(D) - PPL_{\hat{M}(D)}|$$

(Absolute difference)

$$\text{Perplexity Ratio} = \frac{PPL_{\hat{M}}(D)}{PPL_M(D)}$$

(Relative difference)

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Functional Similarity Metrics: Measure Copy Success

d) Perplexity Similarity: A continuous measure of functional extraction success by comparing cross-perplexity between models.

Then the Perplexity Similarity between two models is measured using either absolute difference or relative ratio:

$$\text{Perplexity Gap} = |PPL_M(D) - PPL_{\hat{M}(D)}|$$

(Absolute difference)

$$\text{Perplexity Ratio} = \frac{PPL_{\hat{M}}(D)}{PPL_M(D)}$$

(Relative difference)

- A lower gap or a ratio close to 1 means the extracted model behaves similarly to the target model in how it predicts next tokens.
- A high gap or large deviation in ratio implies the models diverge significantly in their internal distributional behavior.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

(2) How much sensitive data is exposed?



- a) **Training Data Extraction Rate:** % of training data recovered.
- b) **Precision & Recall:** Accuracy and completeness for structured data.
- c) **PII Exposure Rate:** Sensitive user/private info leakage.
- d) **Prompt Recovery Accuracy:** Can system prompts be reconstructed?

Evaluation Measures

Background & Motivation	Taxonomy of LLM MEA	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Evaluation Metrics for Model Extraction Attacks & Defenses

Data Recovery Metrics: Quantifying Info Leakage

a) **Training Data Extraction Rate:** % of training data recovered.

$$TDER = \frac{|\hat{D} \cap D|}{D} \times 100\%$$

D : Set of all training data samples.

\hat{D} : Set of samples recovered (extracted) by the adversary.

$|\hat{D} \cap D|$: Number of correctly extracted training samples.

Evaluation Measures

Background & Motivation	Taxonomy of LLM MEA	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Evaluation Metrics for Model Extraction Attacks & Defenses

Data Recovery Metrics: Quantifying Info Leakage

a) **Training Data Extraction Rate:** % of training data recovered.

Suppose:

The model was trained on $|D|=100,000$ sentences.

The attacker extracts a set \hat{D} containing 2,000 sentences. Among them, 850 sentences match the original training set exactly.

$$TDER = \frac{850}{100,000} \times 100\% = 0.85\%$$

A TDER of 0.85% means 0.85% of the training data was directly leaked, which is a potentially concern if the leaked data includes sensitive or proprietary information.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Data Recovery Metrics: Quantifying Info Leakage

b) Precision & Recall:

- **Precision** measures the proportion of correctly extracted items out of all extracted items.
- **Recall** measures the proportion of correctly extracted items out of all the true sensitive items.

$$\text{Precision} = \frac{|\hat{D} \cap D|}{\hat{D}}$$

$$\text{Recall} = \frac{|\hat{D} \cap D|}{|D|}$$

D : Set of ground-truth sensitive or structured data in the training set;

\hat{D} : Set of data extracted by the attacker;

$\hat{D} \cap D$: Correctly extracted data.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Data Recovery Metrics: Quantifying Info Leakage

b) Precision & Recall:

Suppose:

- Ground-truth sensitive data D includes 200 known email address used during LLM training.
- An attacker extracts a total of 50 email addresses $|\hat{D}| = 50$.
- Out of those, 30 match the original training emails $|\hat{D} \cap D| = 30$.

$$\text{Precision} = \frac{30}{50} = 0.6$$

$$\text{Recall} = \frac{30}{200} = 0.15$$

The attacker is reasonably accurate (60% of their outputs are valid), but has low coverage (only found 15% of all sensitive emails).

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Data Recovery Metrics: Quantifying Info Leakage

c) PII (Personally Identifiable Information) Exposure Rate:

PII quantifies how much sensitive personal information (such as names, email addresses, phone numbers, or social security numbers) has been leaked or reconstructed from an LLM through model extraction attacks.

$$\text{PII Exposure Rate} = \frac{\hat{P} \cap P}{|P|}$$

P: The total set of PII elements present in the model's original training data.

\hat{P} : The set of PII elements extracted by the adversary.

$\hat{P} \cap P$: The number of PII items that were both in the training set and successfully extracted by the attacker.

Evaluation Measures

Background & Motivation	Taxonomy of LLM MEA	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Evaluation Metrics for Model Extraction Attacks & Defenses

Data Recovery Metrics: Quantifying Info Leakage

c) PII (Personally Identifiable Information) Exposure Rate:

Suppose:

- The training dataset contains **1000 pieces** of real PII (e.g., user name, phone numbers, email address, etc), so that $|P|=1000$.
- An attacker successfully extracts 120 strings that resembles PII, with **45 strings match exactly with real training data**, so $|\hat{P} \cap P| = 45$.

$$\text{PII Exposure Rate} = \frac{45}{1000} = 4.5\%$$

4.5% of the original PII from training data has been exposed. Even small exposure rates can have serious privacy implications depending on the nature of the data.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Data Recovery Metrics: Quantifying Info Leakage

d) Prompt Recovery Accuracy: It quantifies how accurately an attacker can reconstruct the original system prompts or instructions used to guide a model's behavior. These prompts often encode sensitive logic, task instructions, or safety constraints.

$$\text{Prompt Recovery Accuracy} = \frac{|\hat{S} \cap S|}{|S|}$$

S : The set of original system prompts or instructions embedded in the target model.

\hat{S} : The set of prompts recovered or reconstructed by the attacker.

$\hat{S} \cap S$: The correctly recovered prompts that match the true underlying instructions.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Data Recovery Metrics: Quantifying Info Leakage

d) Prompt Recovery Accuracy:

Suppose:

- The target model internally uses 10 system prompts for task-specific control (e.g., "Be concise", "Avoid political topics, etc"), so $|S|=10$.
 - After performing a model inversion attack, the adversary reconstructs 6 system prompts, 4 of which are correct and match original ones, i.e., $|\hat{S} \cap S| = 4$
- Then,

$$\text{Prompt Recovery Accuracy} = \frac{4}{10} = 40\%$$

The attacker successfully recovered 40% of the system instructions.
This suggests a partial but significant breach of the model's design or internal safety logic.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

(3) How stealthy & cost-effective is the attack?



- a) The **Effectiveness** of common defense mechanisms on **Functionality Extraction**
- b) The **Effectiveness** of common defense mechanisms on **Training Data Extraction**
- c) The **Effectiveness** of common defense mechanisms on **Prompt-Targeted Attacks**
- d) **Defense Utility:** Preserving Model Value

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Defense Effectiveness Overview

Table: Defense Mechanisms vs. Attack Types

Defense Mechanism	Functionality Extraction			Training Data Extraction		Prompt-targeted Attacks	
	API-based KD	Direct API Querying	Parameter Recovery	Prompt-targeted Recovery	Private Text Reconstruction	Prompt Stealing	Prompt Reconstruction
Architectural Defense [1]	High	Medium	High	Low	Low	Minimal	Minimal
Output Control [2]	High	High	Low	Medium	Medium	Low	Low
Training Data Security [3]	Low	Minimal	Minimal	High	High	Minimal	Minimal
Output Sanitization [4]	Low	Low	Minimal	High	High	Low	Low
Prompt Protection [5]	Minimal	Low	Minimal	Minimal	Minimal	High	High
Query Monitoring [6]	Medium	High	Low	Medium	Medium	Medium	Medium

Effectiveness Levels: High (dark green) - Highly effective; Medium (light green) - Moderately effective; Low (yellow) - Limited effectiveness; Minimal (gray) - Minimal or no effectiveness.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Defense Effectiveness Overview

Table: Defense Mechanisms vs. Attack Types

Defense Mechanism	Functionality Extraction			Training Data Extraction		Prompt-targeted Attacks	
	API-based KD	Direct API Querying	Parameter Recovery	Prompt-targeted Recovery	Private Text Reconstruction	Prompt Stealing	Prompt Reconstruction
Architectural Defense [1]	High	Medium	High	Low	Low	Minimal	Minimal
Output Control [2]	High	High	Low	Medium	Medium	Low	Low
Training Data Security [3]	Low	Minimal	Minimal	High	High	Minimal	Minimal
Output Sanitization [4]	Low	Low	Minimal	High	High	Low	Low
Prompt Protection [5]	Minimal	Low	Minimal	Minimal	Minimal	High	High
Query Monitoring [6]	Medium	High	Low	Medium	Medium	Medium	Medium

Effectiveness Levels: High (dark green) - Highly effective; Medium (light green) - Moderately effective; Low (yellow) - Limited effectiveness; Minimal (gray) - Minimal or no effectiveness.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Defense Effectiveness Overview


Table: Defense Mechanisms vs. Attack Types

Defense Mechanism	Functionality Extraction			Training Data Extraction		Prompt-targeted Attacks	
	API-based KD	Direct API Querying	Parameter Recovery	Prompt-targeted Recovery	Private Text Reconstruction	Prompt Stealing	Prompt Reconstruction
Architectural Defense [1]	High	Medium	High	Low	Low	Minimal	Minimal
Output Control [2]	High	High	Low	Medium	Medium	Low	Low
Training Data Security [3]	Low	Minimal	Minimal	High	High	Minimal	Minimal
Output Sanitization [4]	Low	Low	Minimal	High	High	Low	Low
Prompt Protection [5]	Minimal	Low	Minimal	Minimal	Minimal	High	High
Query Monitoring [6]	Medium	High	Low	Medium	Medium	Medium	Medium


Effectiveness Levels: High (dark green) - Highly effective; Medium (light green) - Moderately effective; Low (yellow) - Limited effectiveness; Minimal (gray) - Minimal or no effectiveness.

Evaluation Measures


Evaluation Metrics for Model Extraction Attacks & Defenses Defense Utility: Preserving Model Value



(1) Performance Preservation:
Minimal impact on intended tasks.



(3) Computation Overhead:
Extra resource cost.



(2) Response Quality:
Maintains generation fluency.

True negative Predicted negative Actual negative	False positive Predicted positive Actual negative
False negative Predicted negative Actual positive	True positive Predicted positive Actual positive

(4) False Positive Rate:
Legitimate queries wrongly blocked.

Evaluation Measures

Evaluation Metrics for Model Extraction Attacks & Defenses

Open Challenges in Evaluation


1) No single metric fits all attack/defense types.



2) Balancing security and usability is hard.



3) Evaluations often empirical, need formal benchmarks.



Part 5: Case Studies & Real-World Scenarios

Case Studies & Real-World Scenarios

Case1: Model Leeching: An Extraction Attack Targeting LLMs

Extracting ChatGPT-3.5-Turbo with just \$50 API cost?



Key findings:

- 73% answer similarity (Exact Match)
- F1 score up to 87%
- Extracted model enables new attacks on LLMs

[1] Birch, Lewis, et al. "Model leeching: An extraction attack targeting llms." *arXiv preprint arXiv:2309.10544* (2023).

Part 5: Case Studies & Real-World Scenarios


Case1: Model Leeching: An Extraction Attack Targeting LLMs

Black-box extraction: Only need public API access, no model details required

Extracting ChatGPT-3.5-Turbo with just \$50 API cost?



Attack Pipeline





[1] Birch, Lewis, et al. "Model leeching: An extraction attack targeting llms." *arXiv preprint arXiv:2309.10544* (2023).

Part 5: Case Studies & Real-World Scenarios

Case1: Model Leeching: An Extraction Attack Targeting LLMs

Extraction Methodology: Prompting, Labeling, and Model Training

Tasks: Question Answering on SQuAD dataset.



- Automated prompt generation to label 83,335 examples.
- Extracted model trained on ChatGPT-labeled data.
- Total cost: \$50.
- Duration: 48hours.

[1] Birch, Lewis, et al. "Model leeching: An extraction attack targeting llms." *arXiv preprint arXiv:2309.10544* (2023).

Part 5: Case Studies & Real-World Scenarios

Case1: Model Leeching: An Extraction Attack Targeting LLMs

Attack Results & Transferability




Fig (a): Model Similarity to ChatGPT-3.5-Turbo. Comparing similarity in correct and incorrect answering of questions relative to ChatGPT-3.5-Turbo.




Fig (b): Baseline and Extracted SQuAD Accuracy. Comparing the baseline and extracted models' performance on the original SQuAD dataset questions and answers.

[1] Birch, Lewis, et al. "Model leeching: An extraction attack targeting llms." *arXiv preprint arXiv:2309.10544* (2023).

Part 5: Case Studies & Real-World Scenarios

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Case1: Model Leeching: An Extraction Attack Targeting LLMs

Why is this important?


- **Low-cost extraction** enables model cloning at scale.
- **Attack transferability:** Stolen models can be used to design new attacks.
- **LLMs served via public APIs are at significant risk.**
- Need for stronger model Intellectual Property protection methods.

[1] Birch, Lewis, et al. "Model leeching: An extraction attack targeting llms." *arXiv preprint arXiv:2309.10544* (2023).

Part 5: Case Studies & Real-World Scenarios

Case 2: Zero-Click LLM Attack: EchoLeak in Microsoft 365 Copilot


- Discovered in Jan 2025 by AimLabs.
- Named EchoLeak, CVE-2025-32711 (CVSS 9.3).
- Allowing silent data exfiltration - NO user interactions required.



[1] Zero-Click AI Vulnerability Exposes Microsoft 365 Copilot Data Without User Interaction. https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios

Case 2: Zero-Click LLM Attack: EchoLeak in Microsoft 365 Copilot



STEP1: Attacker sends a crafted email with hidden prompt injection.

STEP2: Copilot (via RAG) retrieves chunks including malicious payload.

STEP3: Model processes and leaks context data silently.

STEP4: Exfiltration happens automatically via Teams/SharePoint links.

[1] Zero-Click AI Vulnerability Exposes Microsoft 365 Copilot Data Without User Interaction. https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios


Case 2: Zero-Click LLM Attack: EchoLeak in Microsoft 365 Copilot

Key Technical Insights

LLM Scope Violation

What it is: Untrusted email instructions trigger LLM to access privileged data.

Why it works: RAG engine lacks trust segmentation, it treats malicious content as context.



[1] Zero-Click AI Vulnerability Exposes Microsoft 365 Copilot Data Without User Interaction. https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios

Case 2: Zero-Click LLM Attack: EchoLeak in Microsoft 365 Copilot

Key Takeaways & Mitigations

Defenses	Key Takeaways
Patch Copilot (completed June 2025)	Trust boundaries must cover RAG inputs
Restrict external email ingestion (DLP tags)	LLM agents need least-privilege design
Harden prompt and context sanitization (LLM Scope Violation guardrails)	Zero-click attacks are now real threat

[1] Zero-Click AI Vulnerability Exposes Microsoft 365 Copilot Data Without User Interaction. https://thehackernews.com/2025/06/zero-click-ai-vulnerability-exposes.html?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios

Case 3: DeepSeek vs OpenAI: Unintended Model Distillation

OpenAI 'reviewing' allegations that its AI models were used to make DeepSeek

ChatGPT creator warns Chinese startups are 'constantly' using its technology to develop competing products




OpenAI, the developer of ChatGPT, said it knew China-based firms, and others, 'are constantly trying to distil the models of leading US AI companies'. Photograph: GK Images/Alamy

OpenAI has warned that Chinese startups are "constantly" using its technology to develop competing products and said it is "reviewing" allegations that DeepSeek used the ChatGPT maker's AI models to create a rival chatbot.

OpenAI and its partner Microsoft - which has invested \$13bn in the San Francisco-based AI developer - have been investigating whether proprietary technology had been obtained in an unauthorised manner through a technique known as "distillation".

- AI startup DeepSeek reportedly used knowledge distillation on OpenAI's GPT models to build its R1 chatbot.
- Released in January 2025, R1 quickly topped Apple's free app rankings.
- Allegations: model and functionality closely mirror OpenAI's GPT-like capabilities.




[1] OpenAI "reviewing" allegations that its AI models were used to make DeepSeek. https://www.theguardian.com/technology/2025/jan/29/openai-chatgpt-deepseek-china-us-ai-models?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios

Case 3: DeepSeek vs OpenAI: Unintended Model Distillation

How is Distillation Allegedly Performed?

- DeepSeek trained their model using OpenAI API in a black-box manner.
- Technique:



Timeline Highlight:

- 1) Early 2025: R1 released.
- 2) January 2025: OpenAI issues letter alleging unauthorized distillation.

[1] OpenAI "reviewing" allegations that its AI models were used to make DeepSeek. https://www.theguardian.com/technology/2025/jan/29/openai-chatgpt-deepseek-china-us-ai-models?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios

Case 3: DeepSeek vs OpenAI: Unintended Model Distillation

OpenAI & Government Response

OpenAI's Stance:

- (1) Investigating “indications” of unauthorized distillation from GPT.
- (2) Reported evidence and collaborating with US government.

Regulatory Impact:

- (1) US Navy banned DeepSeek usage.
- (2) Added to US tech scrutiny amid rising security concerns.



[1] OpenAI “reviewing” allegations that its AI models were used to make DeepSeek. https://www.theguardian.com/technology/2025/jan/29/openai-chatgpt-deepeek-china-us-ai-models?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Case 3: DeepSeek vs OpenAI: Unintended Model Distillation

Why This Matters?

- **Intellectual Property Theft Risk**
- **Model Development Cost**



- **Market disruption**



[1] OpenAI "reviewing" allegations that its AI models were used to make DeepSeek. https://www.theguardian.com/technology/2025/jan/29/openai-chatgpt-deepseek-china-us-ai-models?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios

Case 3: DeepSeek vs OpenAI: Unintended Model Distillation

Key Takeaways & Mitigations

Lessons Learned	Defenses
Distillation enables IP leakage through black-box API	Rate limits, API monitoring
Market value of covert knowledge transfer is high	Require usage licenses for downstream models
Open-source vs proprietary tension intensifies global race	Regulatory guidelines on model derivation

[1] OpenAI "reviewing" allegations that its AI models were used to make DeepSeek. https://www.theguardian.com/technology/2025/jan/29/openai-chatgpt-deepseek-china-us-ai-models?utm_source=chatgpt.com

Part 5: Case Studies & Real-World Scenarios

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Case 4: Policy Puppetry: Universal Prompt Injection Bypass



- 1) Reported by HiddenAYER Company.
- 2) They discovered attack strategies to bypass guardrails across major LLMs including **GPT-4**, **Claude**, **Gemini**, **Copilot**, **Llama**, **DeepSeek**, etc.
- 3) Enables system-level prompt and harmful content extraction.



[1] Novel Universal Bypass for All Major LLMs -- The Policy Puppetry Prompt Injection Technique: https://hiddenlayer.com/innovation-hub/novel-universal-bypass-for-all-major-langs/?utm_source=chatgpt.com#The-Policy-Puppetry-Attack

Part 5: Case Studies & Real-World Scenarios

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Case 4: Policy Puppetry: Universal Prompt Injection Bypass

Attack Mechanism: How Policy Puppetry Works?

Technique:

Craft malicious prompt formatted as policy file (e.g., XML, JSON)

Effect:

- 1) Overrides model's refusal blocks & alignment.
- 2) Works across different architectures and instruction hierarchies.

[1] Novel Universal Bypass for All Major LLMs -- The Policy Puppetry Prompt Injection Technique: https://hiddenlayer.com/innovation-hub/novel-universal-bypass-for-all-major-langs/?utm_source=chatgpt.com#The-Policy-Puppetry-Attack

Part 5: Case Studies & Real-World Scenarios

Case 4: Policy Puppetry: Universal Prompt Injection Bypass

Attack Effectiveness.

Provider	Model	Effective
OpenAI	ChatGPT 4o-mini	Yes
OpenAI	ChatGPT 4o	Yes
OpenAI	ChatGPT 4.5 Preview	Yes
OpenAI	ChatGPT 4.1	Yes
OpenAI	ChatGPT o1	Yes (with minor adjustments)
OpenAI	ChatGPT o3-mini	Yes (with minor adjustments)
Anthropic	Claude 3.5 Sonnet	Yes
Anthropic	Claude 3.7 Sonnet	Yes
Google	Gemini 1.5 Flash	Yes
Google	Gemini 2.0 Flash	Yes
Google	Gemini 2.5 Pro Preview	Yes (with minor adjustments)
Microsoft	Copilot	Yes

Meta	Llama 3.1 70B Instruct Turbo	Yes
Meta	Llama 3.1 405B Instruct Turbo	Yes
Meta	Llama 3.3 70B Instruct Turbo	Yes
Meta	Llama 4 Scout 17B 16E Instruct	Yes
Meta	Llama 4 Maverick 17B 128E Instruct FP8	Yes
DeepSeek	DeepSeek V3	Yes
DeepSeek	DeepSeek R1	Yes
Qwen	Qwen2.5 72B	Yes
Mistral AI	Mixtral 8x22B	Yes

Demonstrated Impact.

- 1) Elicit harmful content: CBRN instructions, violence, self-harm.
- 2) Leak system prompts & internal instructions.
- 3) Works on agentic systems (with tool access).


[1] Novel Universal Bypass for All Major LLMs -- The Policy Puppetry Prompt Injection Technique: https://hiddenlayer.com/innovation-hub/novel-universal-bypass-for-all-major-langs/?utm_source=chatgpt.com#The-Policy-Puppetry-Attack

Part 5: Case Studies & Real-World Scenarios

Case 4: Policy Puppetry: Universal Prompt Injection Bypass

Why it's So Dangerous?

- Model-agnostic: A single prompt works on GPT, Claude, Copilot, Llama, DeepSeek, Qwen, etc.
- Hard to patch: Rooted in training data; RLHF alone ineffective.
- Scale of threat: Zero-day when developed to consumer apps.



[1] Novel Universal Bypass for All Major LLMs -- The Policy Puppetry Prompt Injection Technique: https://hiddenlayer.com/innovation-hub/novel-universal-bypass-for-all-major-langs/?utm_source=chatgpt.com#The-Policy-Puppetry-Attack

Part 5: Case Studies & Real-World Scenarios

Case 4: Policy Puppetry -- Universal Prompt Injection Bypass

Key Takeaways & Mitigations

Defense	Explanation
Layered Monitoring	Real-time detection of policy-style prompts
Limit Agent Privileges	Avoid unrestricted tool access & minimize context scope
Automated Red-Teaming	Use universal bypass prompts in testing
Incident Playbooks	Prepare responses for jailbreak events

[1] Novel Universal Bypass for All Major LLMs -- The Policy Puppetry Prompt Injection Technique: https://hiddenlayer.com/innovation-hub/novel-universal-bypass-for-all-major-langs/?utm_source=chatgpt.com#The-Policy-Puppetry-Attack

Part 6: Future Directions & Discussions

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------



SECTION OVERVIEW.

- 1) Challenges in LLM Attack.
- 2) Challenges in LLM Defense.
- 3) Roadmap for advancing secure and robust LLMs.

Part 6: Future Directions & Discussions

Challenges in LLM Attack


Limited Model Access & High Cost.

Research gap:

Most attacks in literature use unrealistic unlimited-query assumptions.



**(1) Closed-source Models,
Expensive APIs**



**(2) Unrealistic Unlimited-
Query Assumptions**


Part 6: Future Directions & Discussions

Challenges in LLM Attack


Limited Model Access & High Cost.

Research gap:

Most attacks in literature use unrealistic unlimited-query assumptions.



(1) Closed-source Models,
Expensive APIs



(2) Unrealistic Unlimited-
Query Assumptions


Part 6: Future Directions & Discussions

Challenges in LLM Attack


Limited Model Access & High Cost.

Research gap:

Most attacks in literature use unrealistic unlimited-query assumptions.



(1) Closed-source Models,
Expensive APIs



(2) Unrealistic Unlimited-
Query Assumptions

Future Directions:

Develop query-efficient, stealthy extraction strategies.

Part 6: Future Directions & Discussions


Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Attack

Attack Specificity & Lack of Generalization.

Research Gaps:

- 1) Most extraction attacks exploit isolated model features (e.g., output tokens, logits).
- 2) Attacks rarely scale across architectures or tasks.
- 3) Few studies address **adaptive** or **multi-pronged** extraction.




Part 6: Future Directions & Discussions

Challenges in LLM Attack

Attack Specificity & Lack of Generalization.

Research Gaps:

- 1) Most extraction attacks exploit isolated model features (e.g., output tokens, logits).
- 2) Attacks rarely scale across architectures or tasks.
- 3) Few studies address **adaptive** or **multi-pronged** extraction.



Future Directions:

Combine diverse attack vectors to defeat adaptive defenses.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Attack

Stealth vs. Effectiveness Trade-off.

Research Gaps:

- 1) High-fidelity extraction needs massive queries—risks detection and cost escalation.
- 2) Stealthier attacks often reduce extraction quality.
- 3) Balancing cost, risk, and model fidelity remains unsolved.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Attack

Stealth vs. Effectiveness Trade-off.

Research Gaps:

- 1) High-fidelity extraction needs massive queries—risks detection and cost escalation.
- 2) Stealthier attacks often reduce extraction quality.
- 3) Balancing cost, risk, and model fidelity remains unsolved.

Future Directions:

Leverage active learning, reinforcement learning for optimal query planning.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Defense

Current Defense Limitations.

- 1) Structural defenses (e.g., model watermarking, API filtering) are hard to deploy on production models.
- 2) Output randomization harms utility/accuracy.
- 3) Most defenses lack formal guarantees; mostly evaluated empirically.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Defense

Current Defense Limitations.

- 1) Structural defenses (e.g., model watermarking, API filtering) are hard to deploy on production models.
- 2) Output randomization harms utility/accuracy.
- 3) Most defenses lack formal guarantees; mostly evaluated empirically.

Future Direction:

Research plug-and-play defenses for black-box models

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Defense

Cat-and-Mouse: Arms Race Continues.

Research Gaps:

- 1) Adaptive attackers quickly bypass static defenses.
- 2) Defenses based on output manipulation can often be reverse-engineered.



Part 6: Future Directions & Discussions

Challenges in LLM Defense

Cat-and-Mouse: Arms Race Continues.

Research Gaps:

- 1) Adaptive attackers quickly bypass static defenses.
- 2) Defenses based on output manipulation can often be reverse-engineered.



Future Direction:

Defenses must anticipate adversarial adaptation.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Defense

Need of Formal Security Guarantees.

Research Gaps:

- 1) Most current evaluations are empirical; few offer theoretical security.
- 2) No standardized benchmarks or threat metrics.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Defense

Need of Formal Security Guarantees.

Research Gaps:

- 1) Most current evaluations are empirical; few offer theoretical security.
- 2) No standardized benchmarks or threat metrics.

Future Directions:

- 1) Develop provable defenses (cryptographic, information-theoretic).
- 2) Draw on work from differential privacy, watermarking, and robust learning.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Defense

Defense Applicability & Usability Gaps.

Research Gaps:

- 1) Most defenses require access to model internals or retraining.
- 2) Few methods can retrofit existing deployed APIs.
- 3) Defenses must not hurt model performance or UX.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Challenges in LLM Defense

Defense Applicability & Usability Gaps.

Research Gaps:

- 1) Most defenses require access to model internals or retraining.
- 2) Few methods can retrofit existing deployed APIs.
- 3) Defenses must not hurt model performance or UX.

Future Directions:

Focus on post-deployment, non-invasive methods.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Roadmap for advancing secure and robust LLMs

Expanding Threat & Evaluation Scenarios.

Research Gaps:

- 1) Most research focuses on QA/classification; other tasks (code, multimodal, agentic) are underexplored
- 2) Extraction in federated, on-device, and collaborative LLMs?

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Roadmap for advancing secure and robust LLMs

Expanding Threat & Evaluation Scenarios.

Research Gaps:

- 1) Most research focuses on QA/classification; other tasks (code, multi-modal, agentic) are underexplored
- 2) Extraction in federated, on-device, and collaborative LLMs?

Future Directions:

Build diverse, realistic benchmarks & red-teaming scenarios.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Roadmap for advancing secure and robust LLMs

Vision for Robust LLM Ecosystem.

Long-term Vision: Secure and Trustworthy LLMs

- 1) Industry–academia collaboration for shared threat intelligence.
- 2) Regulation and best practices for LLM APIs.
- 3) Red-teaming, open benchmarks, and public reporting.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Roadmap for advancing secure and robust LLMs

Vision for Robust LLM Ecosystem.

Long-term Vision: Secure and Trustworthy LLMs

- 1) Industry–academia collaboration for shared threat intelligence.
- 2) Regulation and best practices for LLM APIs.
- 3) Red-teaming, open benchmarks, and public reporting.

Future Direction:

Ongoing research is critical for future-proof LLMs.

Part 6: Future Directions & Discussions

Background & Motivation	Taxonomy of Attacks	Defense Techniques	Evaluations	Case Studies	Future Directions
-------------------------	---------------------	--------------------	-------------	--------------	-------------------

Thank you for listening!

Q & A



We welcome your questions!

